




OxML2024: LLM & Diffusion Model

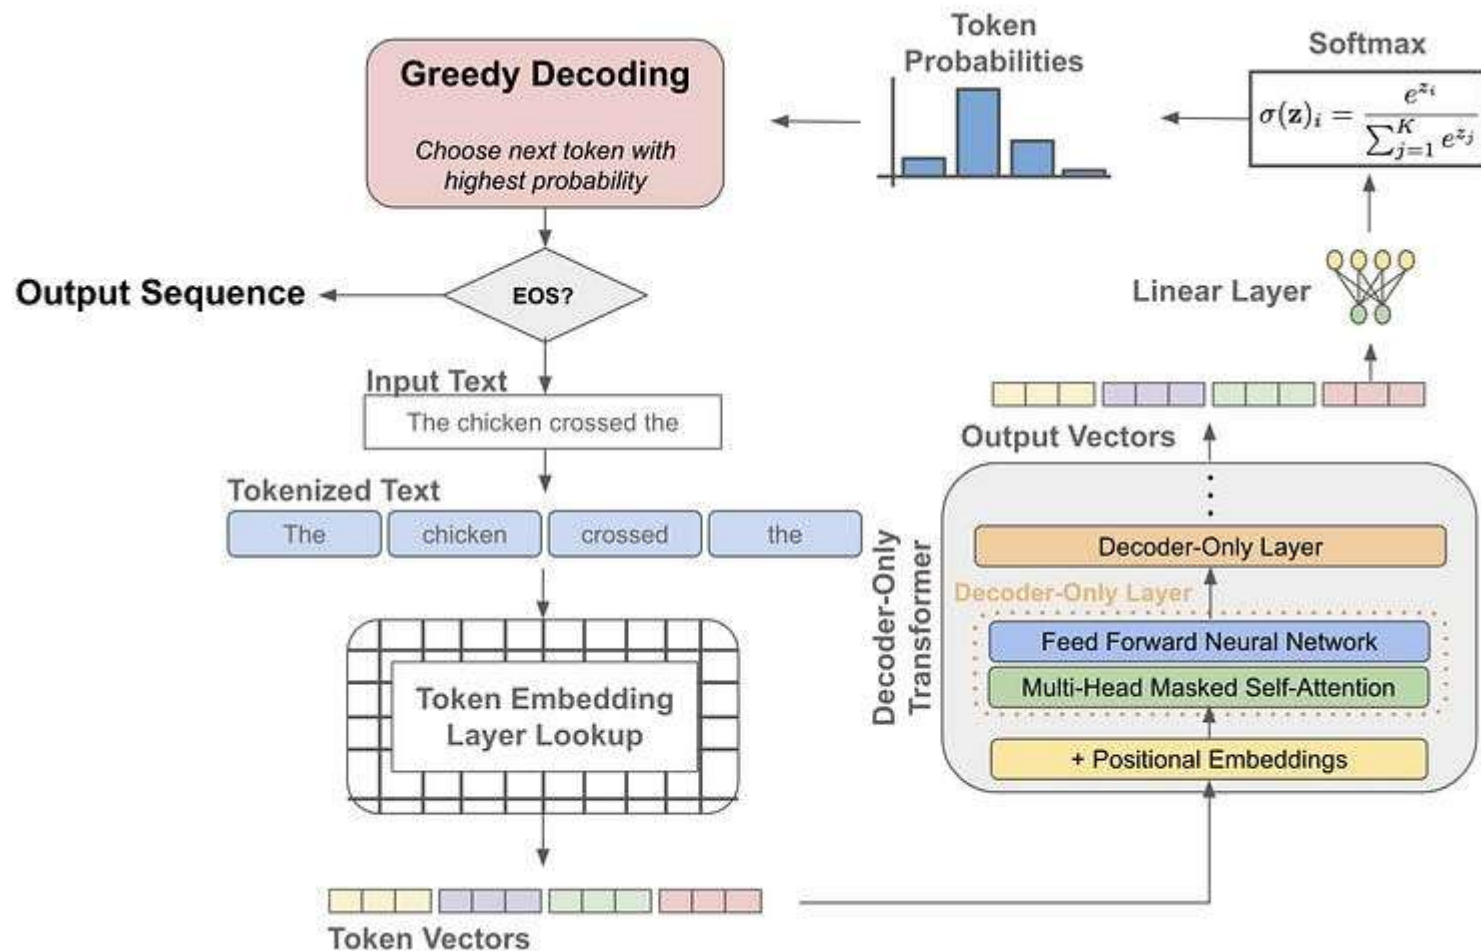
 Wenhan Han
TU Eindhoven

Overview

Colab Notebook: [LLM Link](#) [Diffusion Link](#)

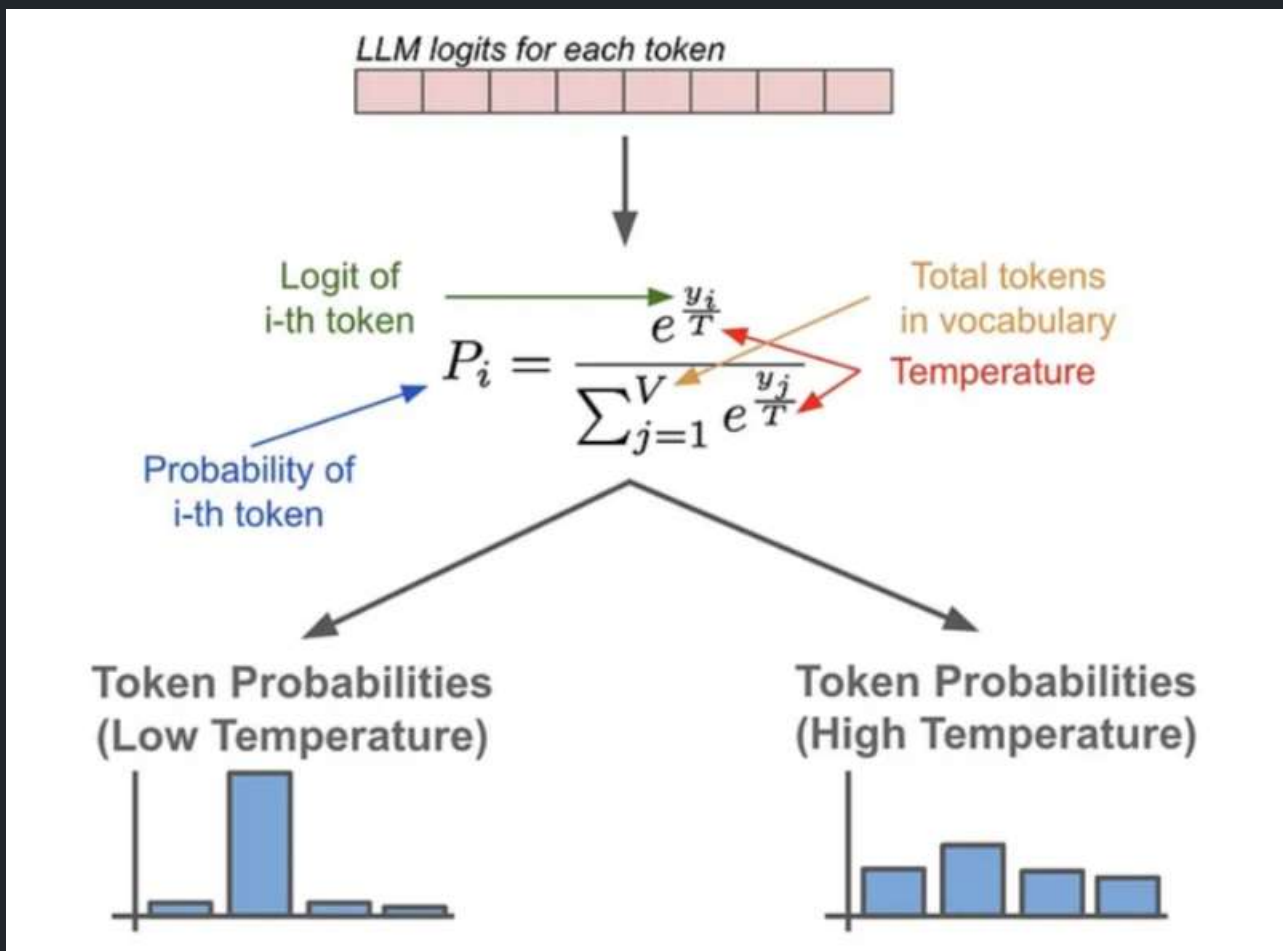
- Access to **LLMs** through APIs.
- Run **LLM** inference in local.
- Prompting techniques.
- Finetune **LLaMA-3** on single GPU.
- Run **Stable Diffusion**.
- Train a mini **Diffusion Model**.

Recall: LLM Inference (Decoder Only)



Recall: LLM Inference (Decoder Only)

Sampling Parameter: Temperature




In the context of a language model, the probability of the next token (word or sub-word unit) is given by a distribution P , which is a function of the model's internal parameters and the input context.

High Temperature: The scaling effectively "flattens" the distribution, making less probable tokens more likely to be sampled.


Low Temperature: The distribution becomes "sharper," concentrating the probability mass on the most likely tokens.

Quantization




The two most common 8-bit quantization techniques are zero-point quantization and absolute maximum (absmax) quantization. Zero-point quantization and absmax quantization map the floating point values into more compact int8 (1 byte) values.

Quantization




Quantization

Handle Outliers



- From the input hidden states, extract the **outliers** (i.e. values that are larger than a certain threshold) by column.
- Perform the matrix multiplication of the **outliers** in FP16 and the **non-outliers** in int8.
- Dequantize the **non-outlier** results and add both **outlier** and **non-outlier** results together to receive the full result in FP16.

LoRA

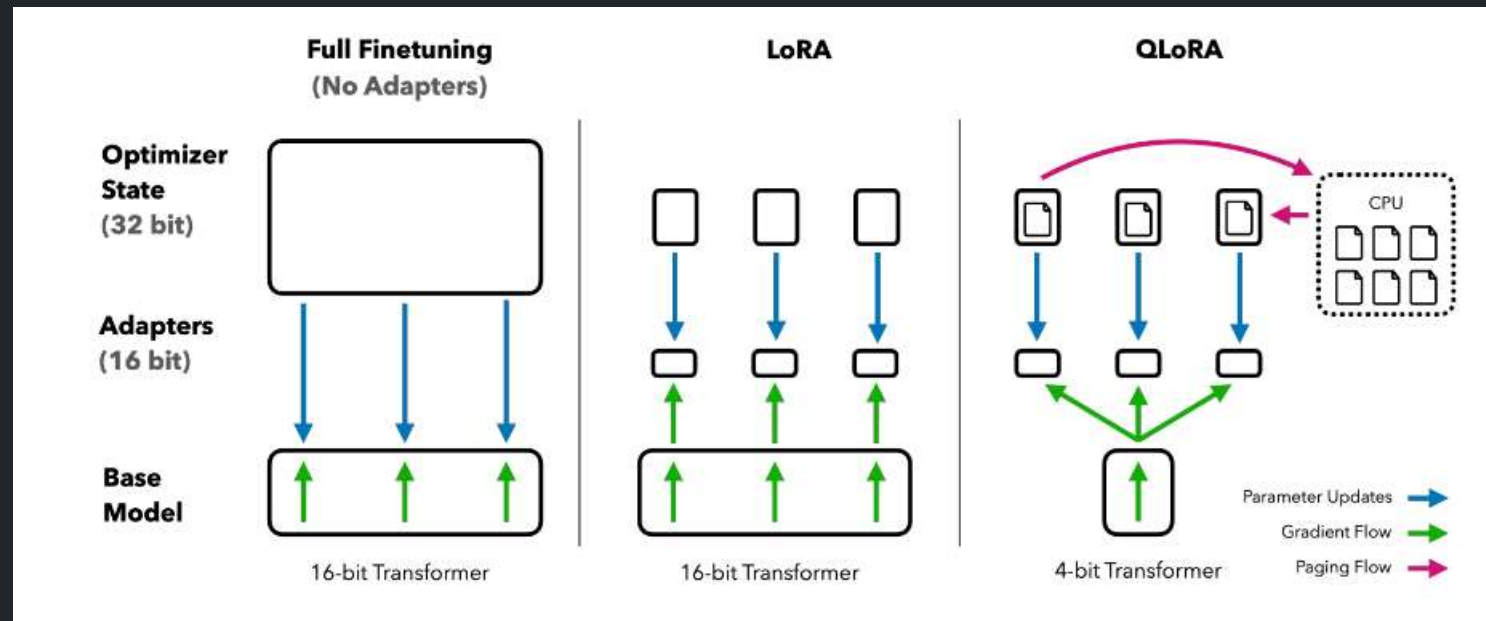


We only train A and B . The finetuned weight becomes

$$W = W_o + \Delta W = W_o + BA$$


W_o denotes the pre-trained parameter weights.

QLoRA



QLoRA is the extended version of LoRA which works by quantizing the precision of the weight parameters in the pre trained LLM to 4-bit precision.

QLoRA



QLoRA introduces multiple innovations designed to reduce memory use without sacrificing performance:

- **4-bit NormalFloat**, an information theoretically optimal quantization data type for normally distributed data that yields better empirical results than 4-bit Integers and 4-bit Floats.
- **Double Quantization**, a method that quantizes the quantization constants, saving an average of about 0.37 bits per parameter (approximately 3 GB for a 65B model).
- **Paged Optimizers**, using NVIDIA unified memory to avoid the gradient checkpointing memory spikes that occur when processing a mini-batch with a long sequence length.

Recall: Diffusion Model




A **Diffusion Model** or Score-based Generative Model perturbs the input data into standard Gaussian noise and learns the reverse process. By sampling from the Gaussian distribution, the **Diffusion Model** can iteratively reconstruct the samples.

Recall: Diffusion Model

To be general, the forward process follows a **Stochastic Differential Equation** (SDE) that has the form as

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$




There is a corresponding reverse SDE

$$d\mathbf{x} = [\mathbf{f}(\mathbf{x}, t) - g^2(t) \nabla_{\mathbf{x}} \log_{p_{\sigma}}(\mathbf{x})]dt + g(t)d\mathbf{w}$$

Score function estimated by a score model

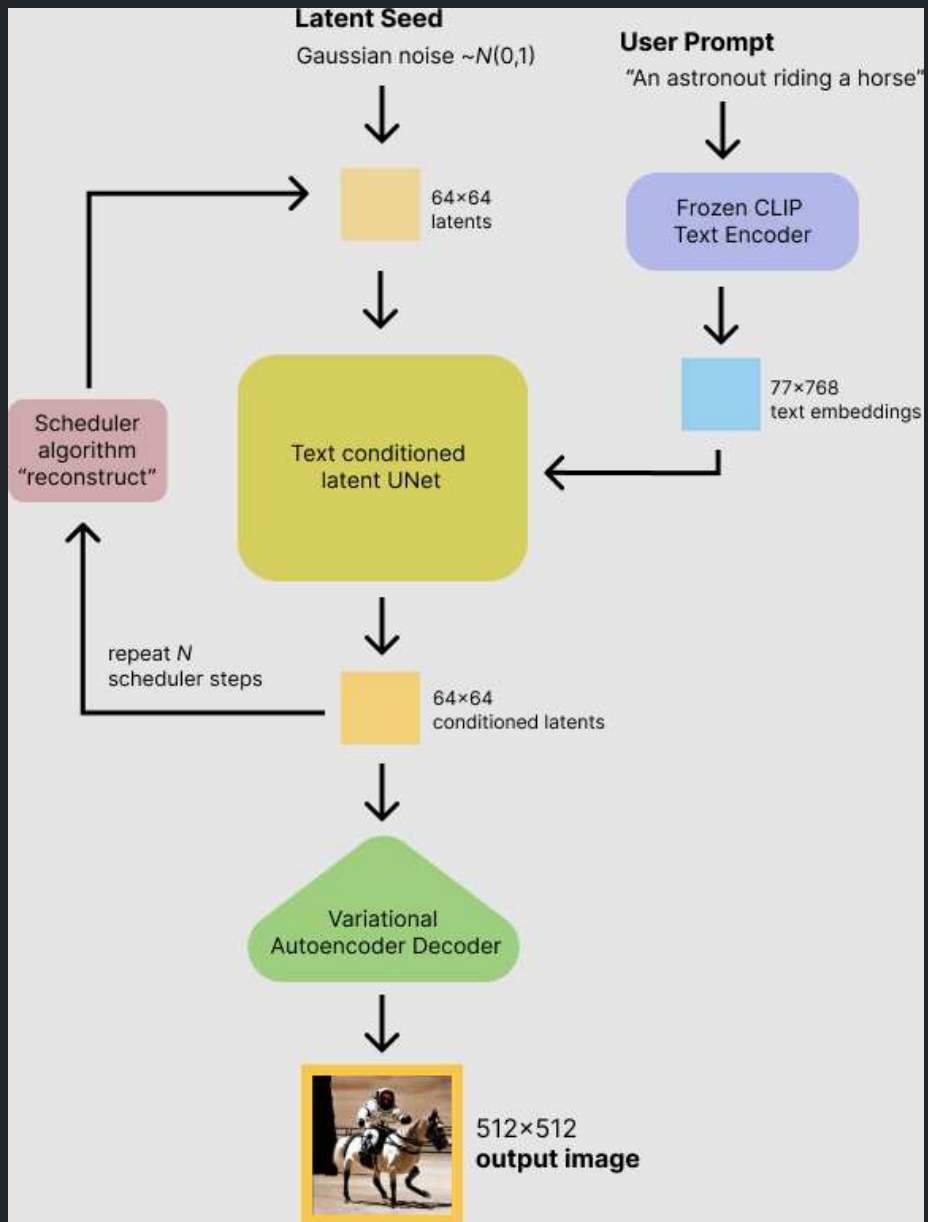
Recall: Diffusion Model



The sampling can follow a simple numerical SDE solver
Euler-Maruyama Method.

$$\begin{array}{|l} \Delta \mathbf{x} = [\mathbf{f}(\mathbf{x}, t) - g^2(t)s_{\theta}(\mathbf{x}, t)]\Delta t + g(t)\sqrt{|\Delta t|}\mathbf{z}_t \\ \mathbf{x} = \mathbf{x} + \Delta \mathbf{x} \\ t = t - \Delta t \end{array}$$

Stable Diffusion



Stable Diffusion is a large text to image diffusion model trained on billions of images.

Thanks